ISTM 6212 - Week 9
Hierarchies, Snapshots

Daniel Chudnov, dchud@gwu.edu

2016-11-08

Agenda

- Schedule check
- Quick review
- Handling Hierarchies
- Snapshots
- Accumulating Snapshots
- Project 02 work session

Schedule check

Quick review

Facts and dimensions

- **Facts** are instances of business processes worthy of measurement

- **Dimensions** are the contexts in which those processes occurred and through which their measurement may be framed

Facts are sparse; dimensions wide

- Facts represent individual events; no records for "all possible events", only what actually happened

- Dimensions represent possible contexts; records for many possible combinations of filter / aggregation attributions

Consistency in dimensions

- Same structure: same attributes, names, types
- Same content: same values, casing, abbreviations
- Queries can account for differences, but early planning and proper ETL can make drilling across easier
- Attributes must match even if tables don't

Affinity in dimensions

- salesperson + territory vs. salesperson + customer
- product + brand + model
- team + player
- is affinity natural, or process/event-based?
- is affinity within one context or several?

Addressing large dimensions

- is it really more than one dimension? split them up.

- extract subtype specifics to new dimensions (e.g. related product types)

- mini-dimensions: split out possibilities

Avoid NULL in dimensions

- Special-case roles for missing or bad data

Used when a fact is supplied with an invalid product_code

Used when a fact arrives prior to dimensional context

PRODUCT

product_key	row_type	product_code	product_name	
0	Invalid	n/a	n/a	
1	Unknown	n/a	n/a	
101	Product	B57330-1	Cardboard Box	
102	Product	B47770-2	Bubble Envelope	

Fig. 6-11

Behavioral dimensions

- Read this closely!
- Use facts to create new dimensions:
  - categorize customers by sales level / frequency
  - categorize products by popularity / seasonality
- A prelude to **feature engineering** in data mining

Handling Hierarchies

Hierarchies appear everywhere

- Many dimensions have inherent hierarchy
- Hierarchies play out in scoping context, summarization
- More than one hierarchy may be equally valid

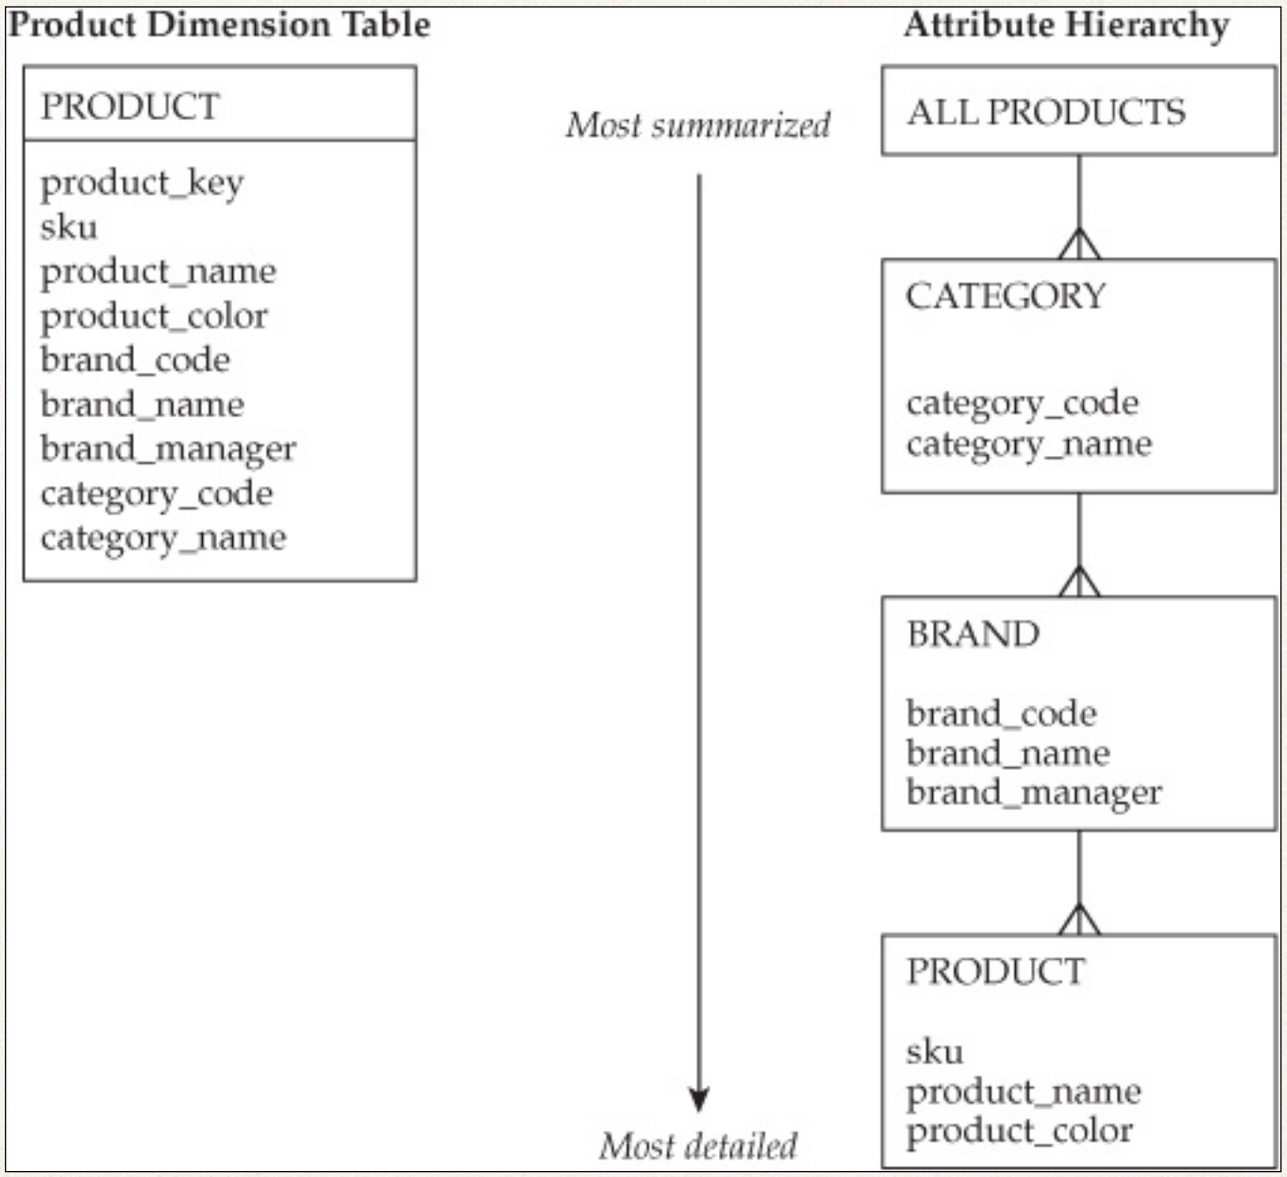

Fig. 7-2

All Products (1) → Categories (25) → Brands (650) → Products (8000)

Instance (recursive) hierarchies

- Organizational structures
  - Departments / sections
  - Owners / subsidiaries
- Employee relationships (e.g. supervisor)
- Product components (parts / assemblies)

Instance hierarchy handling

- Recursive queries
- Flattened dimensions
- Bridge tables

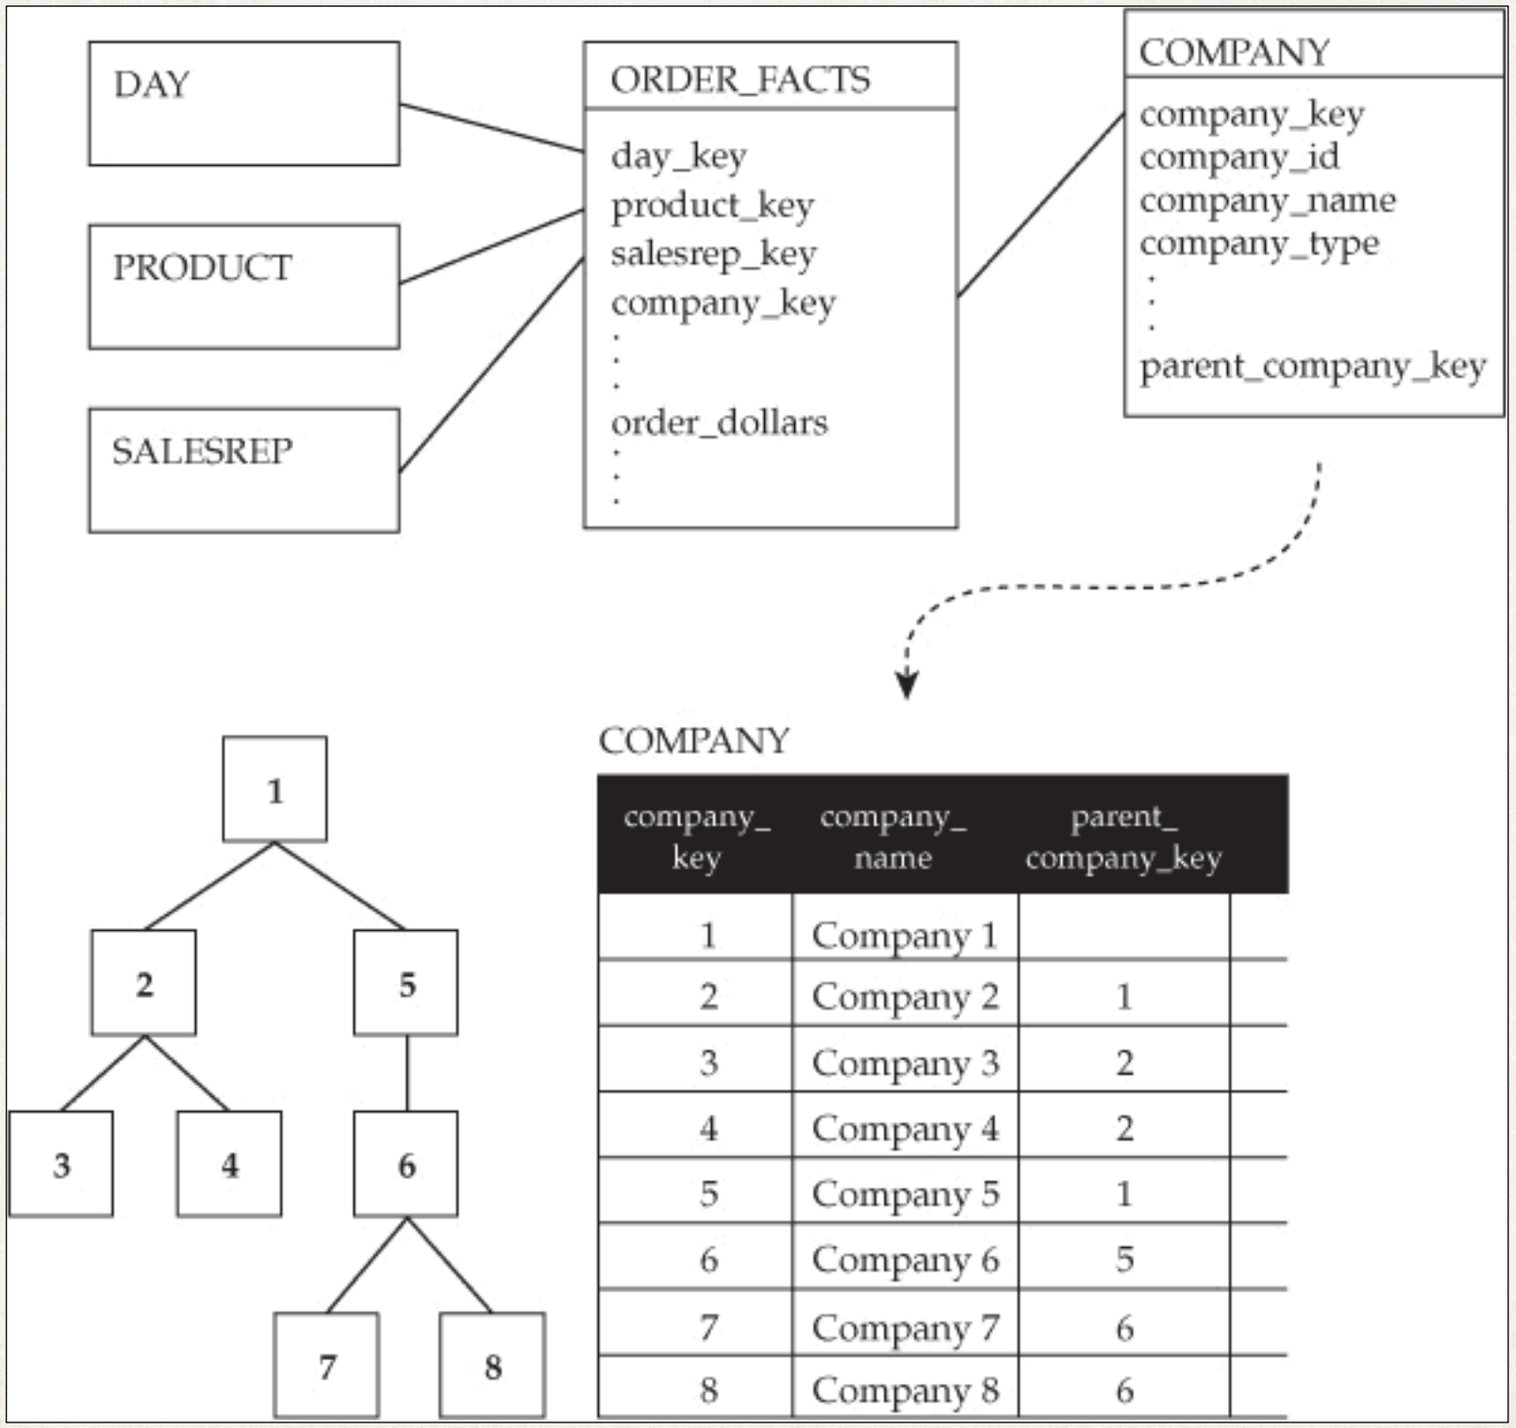

Fig. 10-3

Recursive company dimension structure

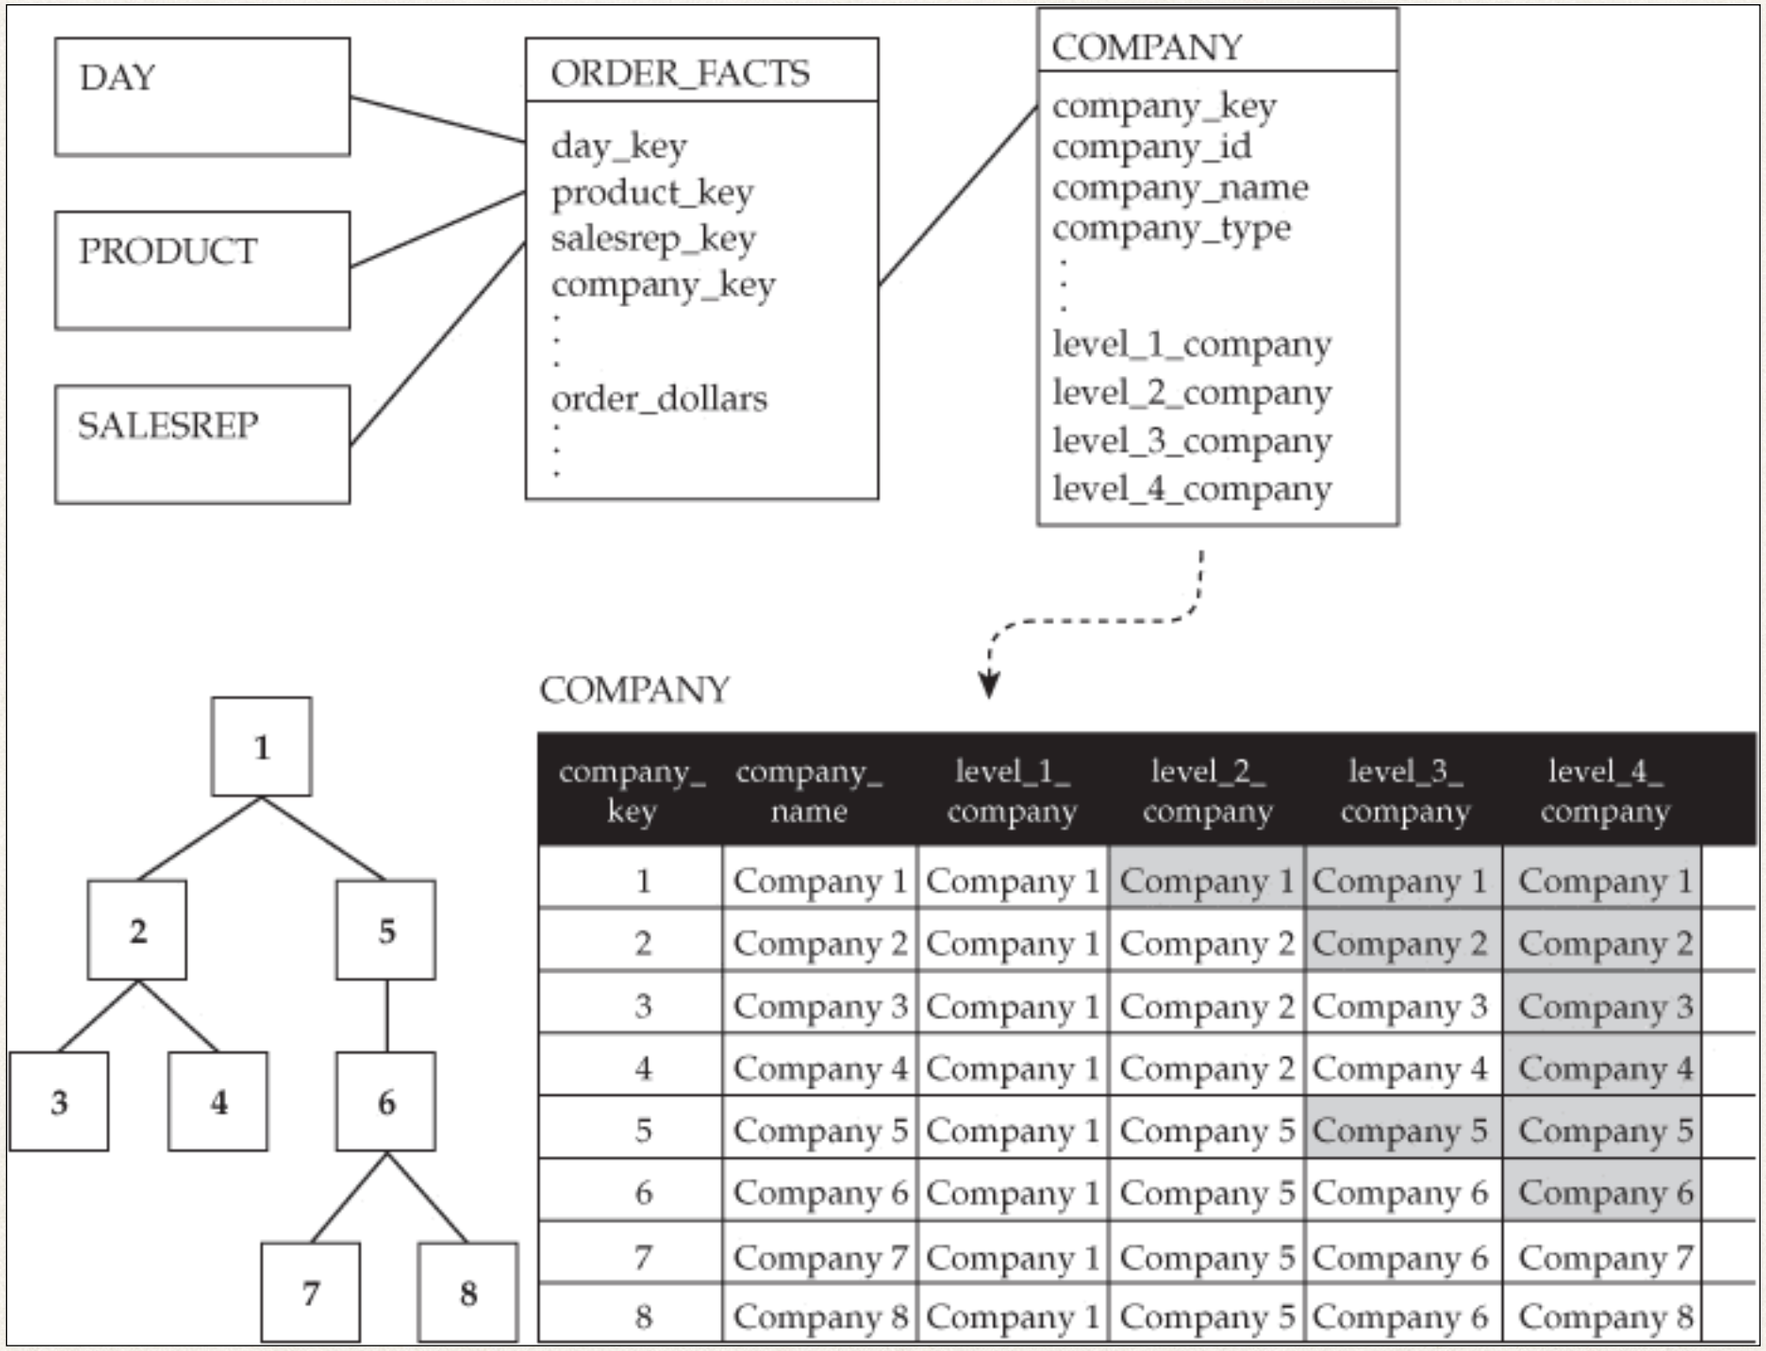

COMPANY

company_ key	company_ name	level_1_ company	level_2_ company	level_3_ company	level_4_ company
1	Company 1	Company 1	Company 1	Company 1	Company 1
2	Company 2	Company 1	Company 2	Company 2	Company 2
3	Company 3	Company 1	Company 2	Company 3	Company 3
4	Company 4	Company 1	Company 2	Company 4	Company 4
5	Company 5	Company 1	Company 5	Company 5	Company 5
6	Company 6	Company 1	Company 5	Company 6	Company 6
7	Company 7	Company 1	Company 5	Company 6	Company 7
8	Company 8	Company 1	Company 5	Company 6	Company 8

Fig. 10-4

Flattened corporate hierarchy dimension

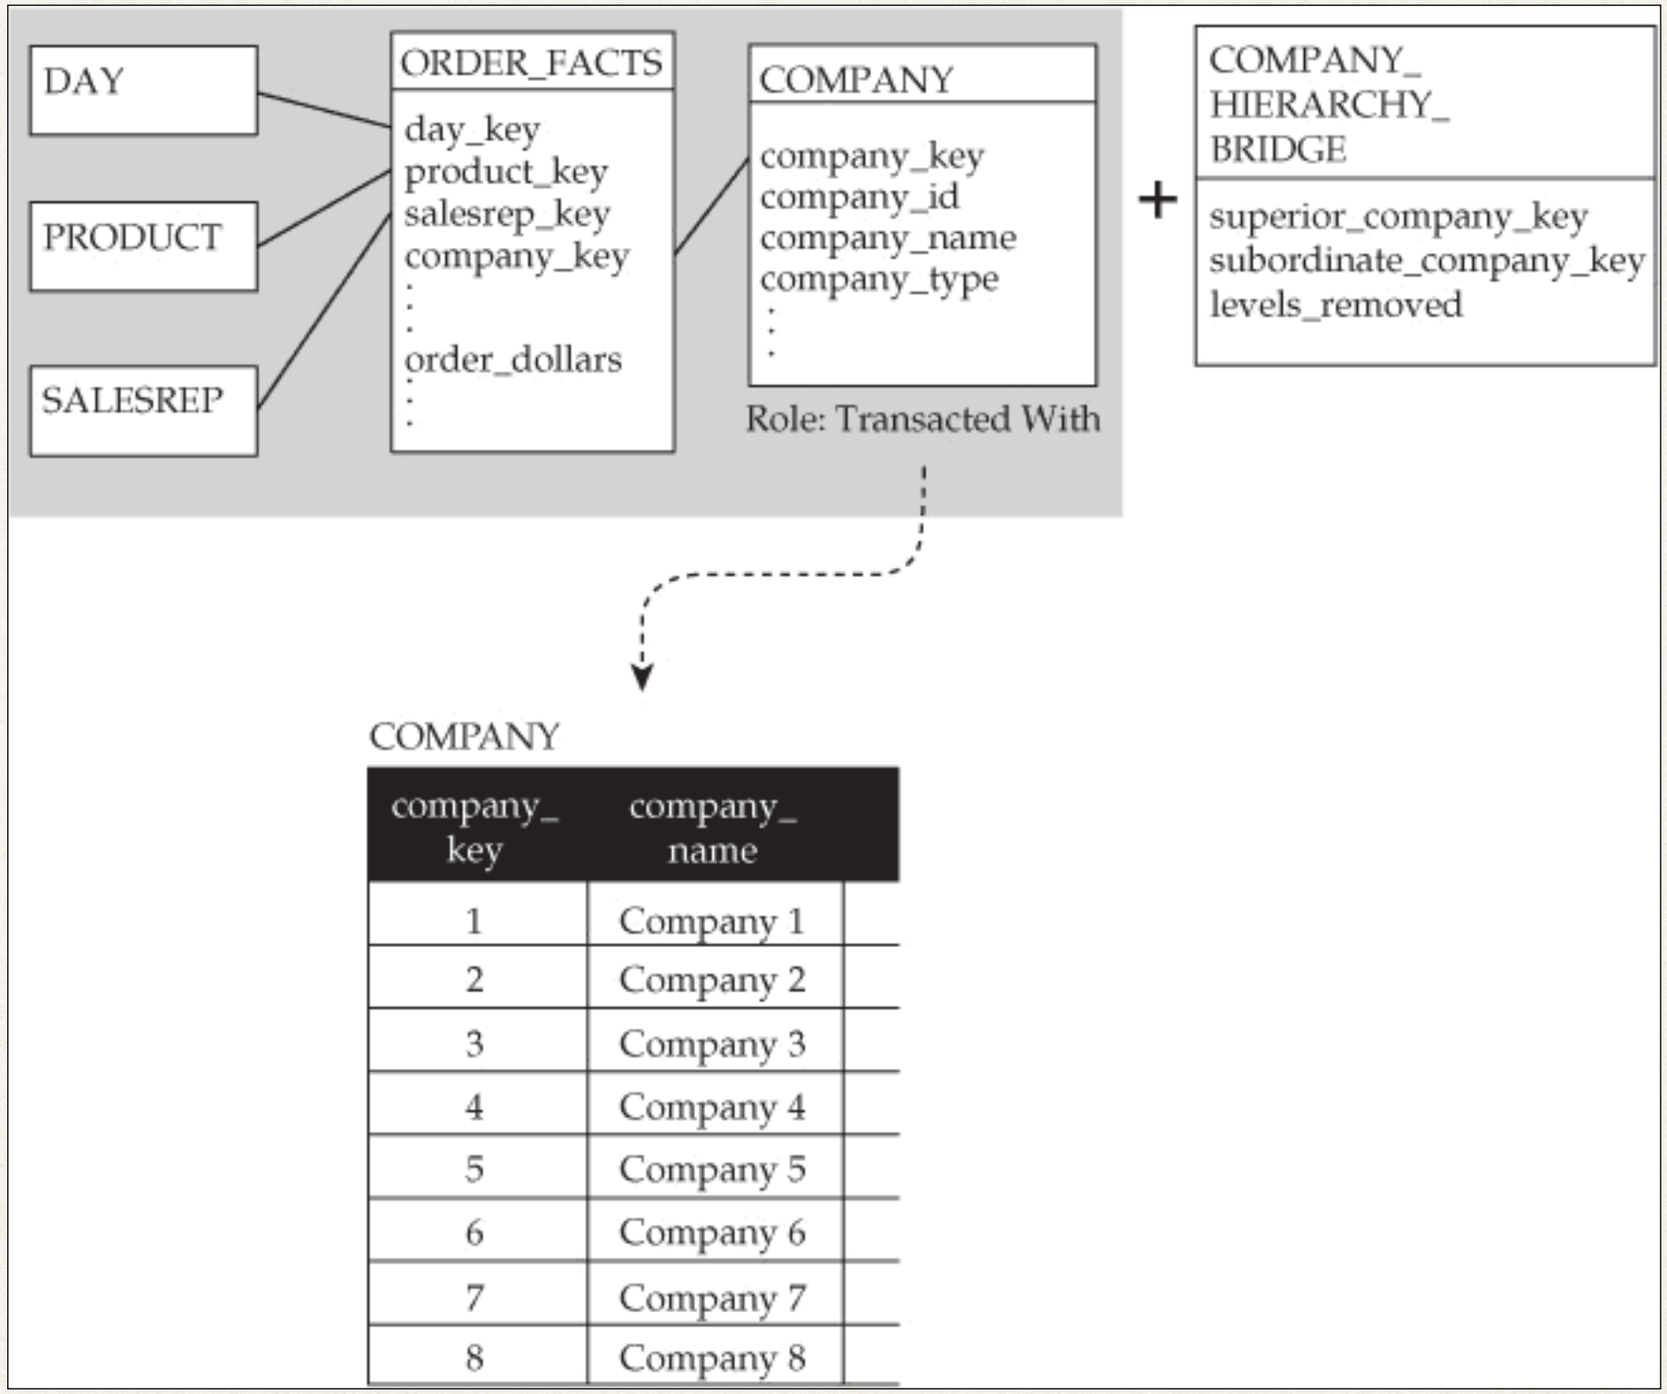

COMPANY

company_ key	company_ name	
1	Company 1	
2	Company 2	
3	Company 3	
4	Company 4	
5	Company 5	
6	Company 6	
7	Company 7	
8	Company 8	

Fig. 10-5

Bridge structure for corporate ownership hierarchy

- these structures occur frequently!
- see *Star Schema,* Chapter 10

Key tactics

- Document attribute hierarchies in dimensions
- Document dimension conformance
- Design to meet analysis needs, not all possible configurations

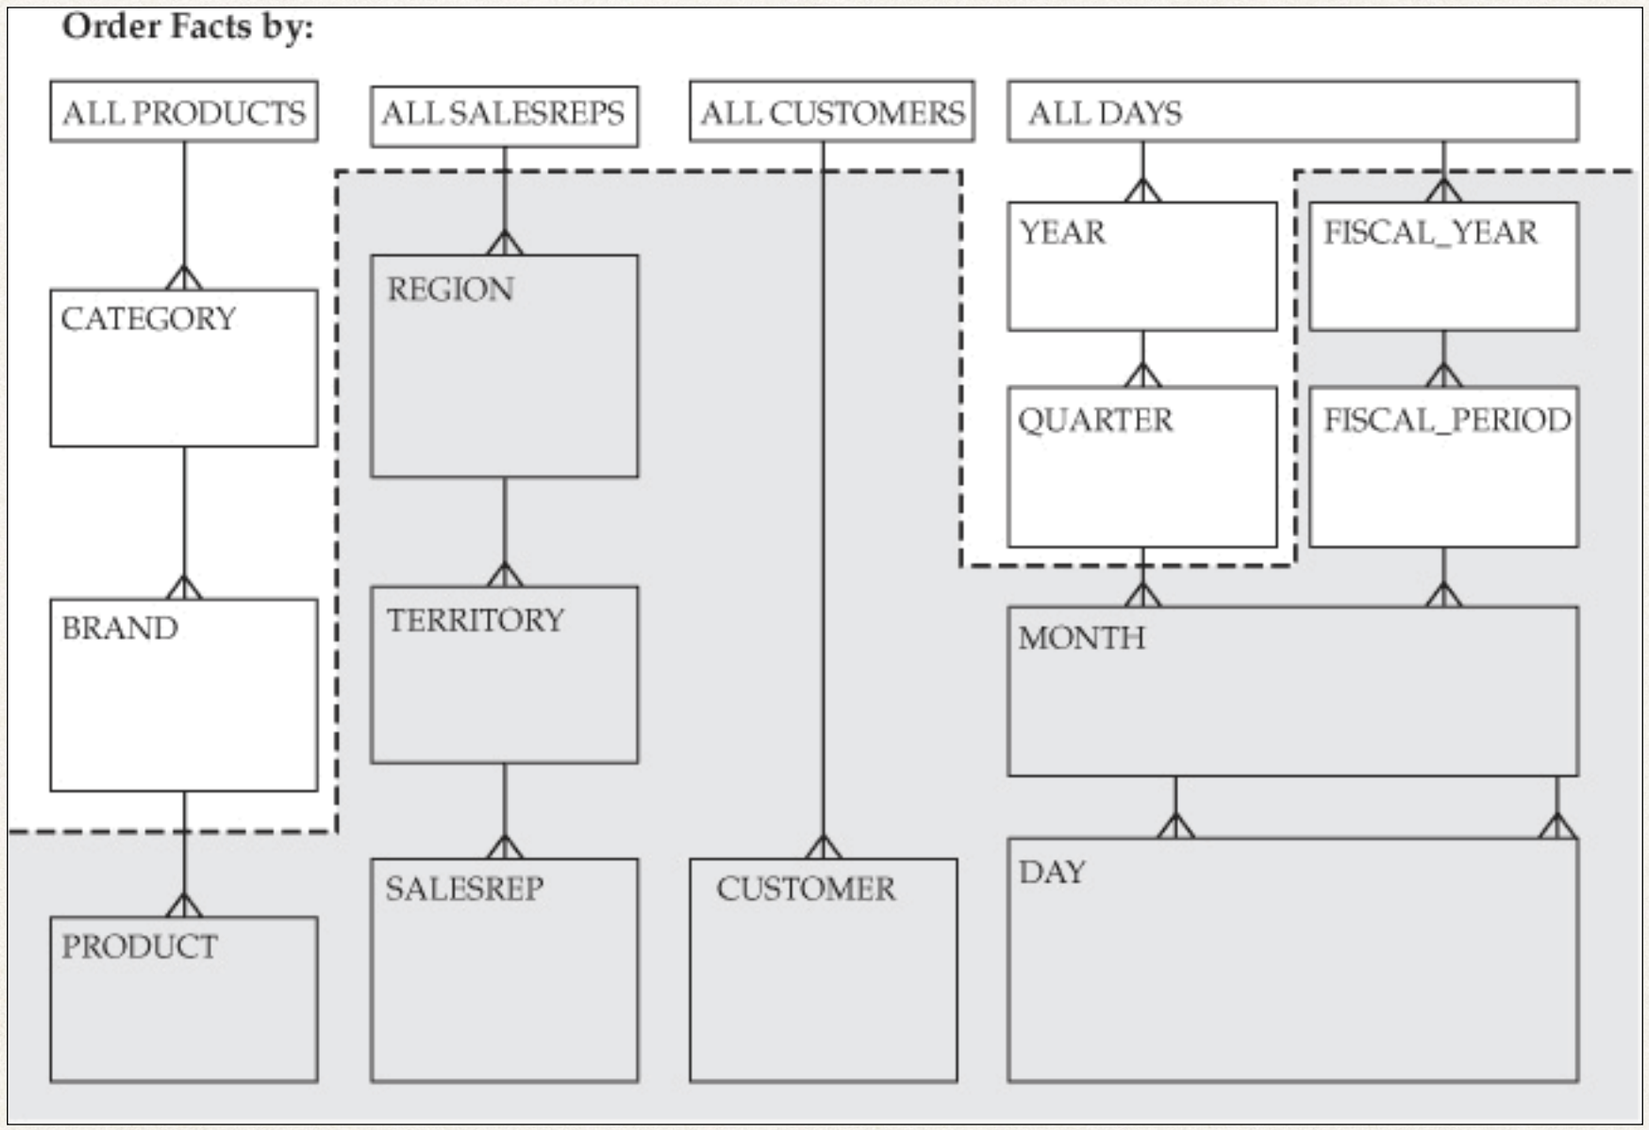

Fig. 7-4

Hierarchy diagram defines aggregate

Quick thoughts on Snowflakes

- Avoid urge to normalize in dimensional models
- Have a good reason:
  - BI products expect or benefit
  - Some database engines optimize for it
  - Multi-value / repeating groups
  - Cleaner ETL / modeling outweighs query / update penalty

Snapshots

So far: Transaction fact tables

- Facts as records of transaction / process / event at a specific granularity

- Facts are additive - sums have useful meaning

- Each record / row represents measurable facts based on what actually occurred

- "Sparse" - records of actualized activity, not theoretical possibilities

Transaction fact tables awkward for:

- Status (inventory levels)
- Sensor readings (changing state)
- Prices or balances (moving averages)
- Not transactions, not additive, not events, not sparse

Bank accounts: txns vs. balances

- Reasonable to study transaction patterns: these fit the transaction fact table model well

- What about demographics or seasonality of balances?

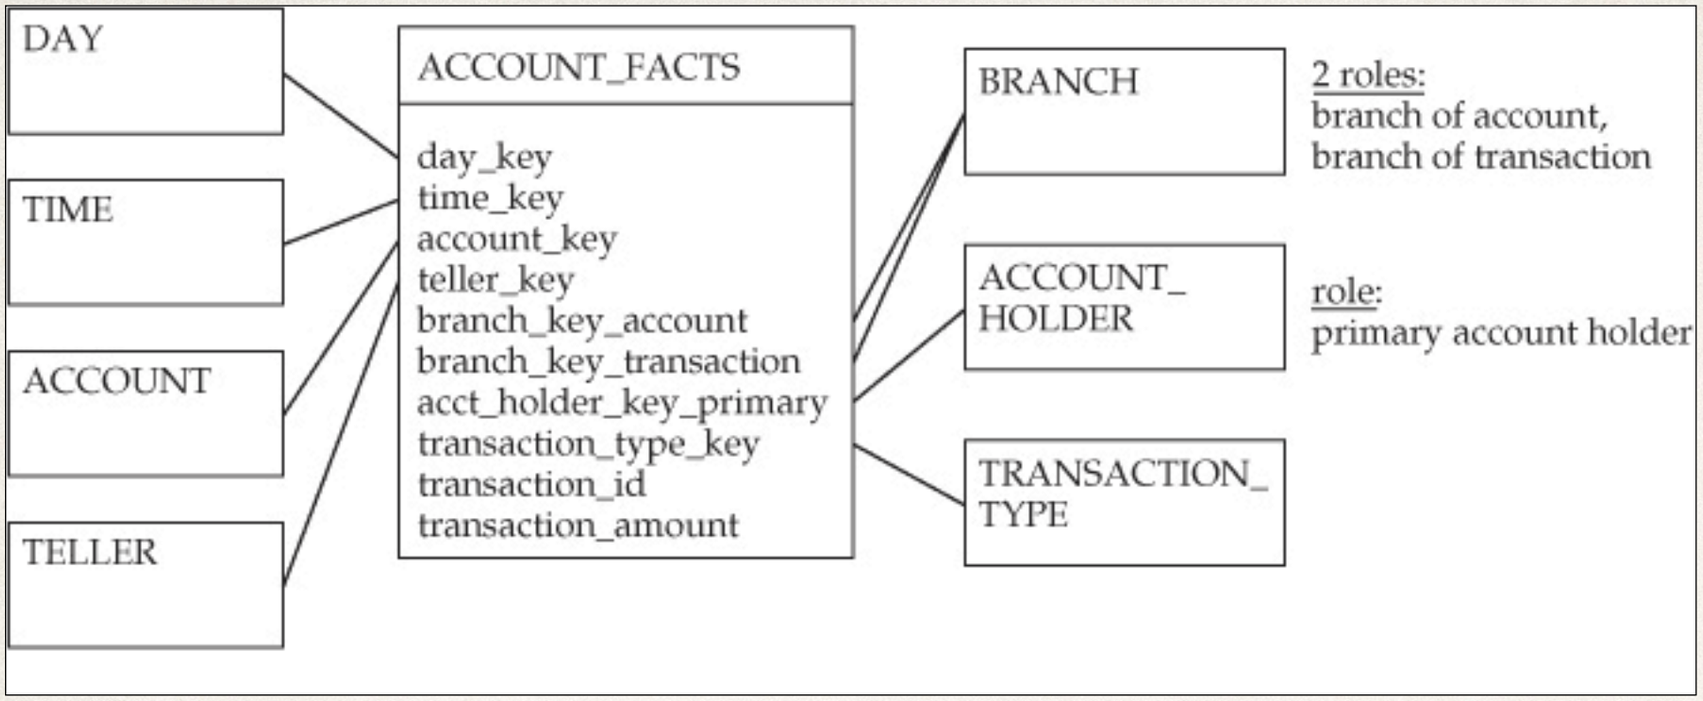

Fig. 11-1

Account transaction fact table design & instances

Account: 7922-3002
Period: 2/1/2009 – 2/14/2009

Granular transaction data stored in star:

Day	Transaction Type	Transaction Amount
2/1/2009	Initial Deposit	2000.00
2/2/2009	Withdrawal	(20.00)
2/3/2009	Check	(35.50)
2/3/2009	Check	(17.02)
2/6/2009	Check	(75.00)
2/6/2009	Deposit	75.00
2/7/2009	Check	(800.00)
2/10/2009	Check	(68.29)
2/14/2009	Withdrawal	(100.00)

Issues with transactions

- Tells us nothing about account balances
- Balances are not additive
- Many days without transactions
- Many days with transactions

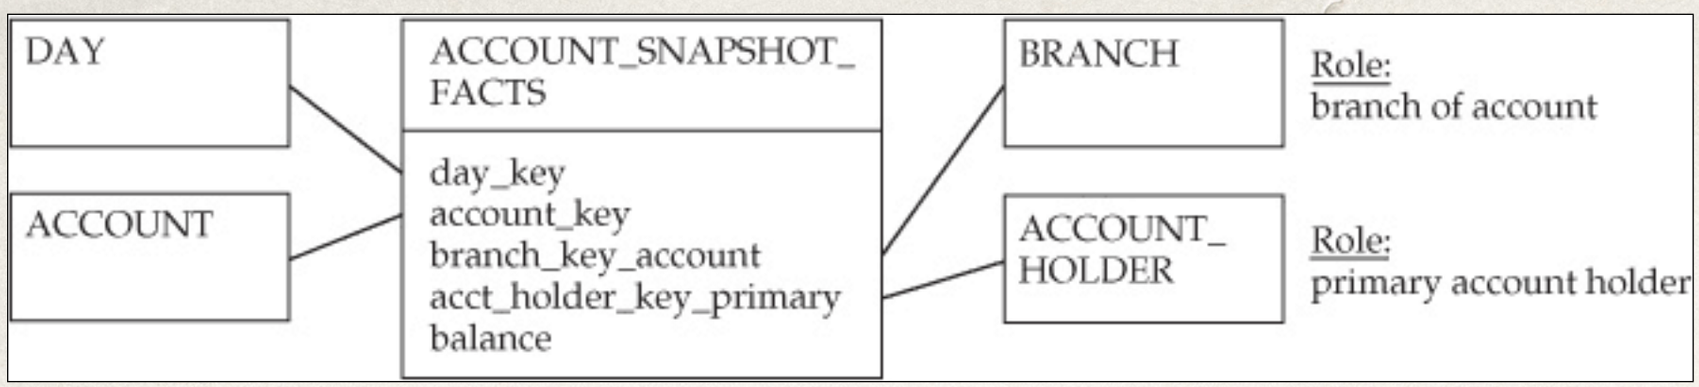

Fig. 11-2

Account
snapshot fact table
design
&
instances

Account: 7922-3002
Period: 2/1/2009 – 2/14/2009

Granular transaction data stored in account_facts:

Day	Transaction Type	Transaction Amount
2/1/2009	Initial Deposit	2000.00
2/2/2009	Withdrawal	(20.00)
2/3/2009	Check	(35.50)
2/3/2009	Check	(17.02)
2/6/2009	Check	(75.00)
2/6/2009	Deposit	75.00
2/7/2009	Check	(800.00)
2/10/2009	Check	(68.29)
2/14/2009	Withdrawal	(100.00)

Periodic status data stored in account_snapshot_facts:

Day	Balance
2/1/2009	2000.00
2/2/2009	1980.00
2/3/2009	1927.48
2/4/2009	1927.48
2/5/2009	1927.48
2/6/2009	1927.48
2/7/2009	1127.48
2/8/2009	1127.48
2/9/2009	1127.48
2/10/2009	1059.19
2/11/2009	1059.19
2/12/2009	1059.19
2/13/2009	1059.19
2/14/2009	959.19

Two transactions occurred on this day

No transactions occurred on these days

Offsetting transactions occurred on this day

Fig. 11-3

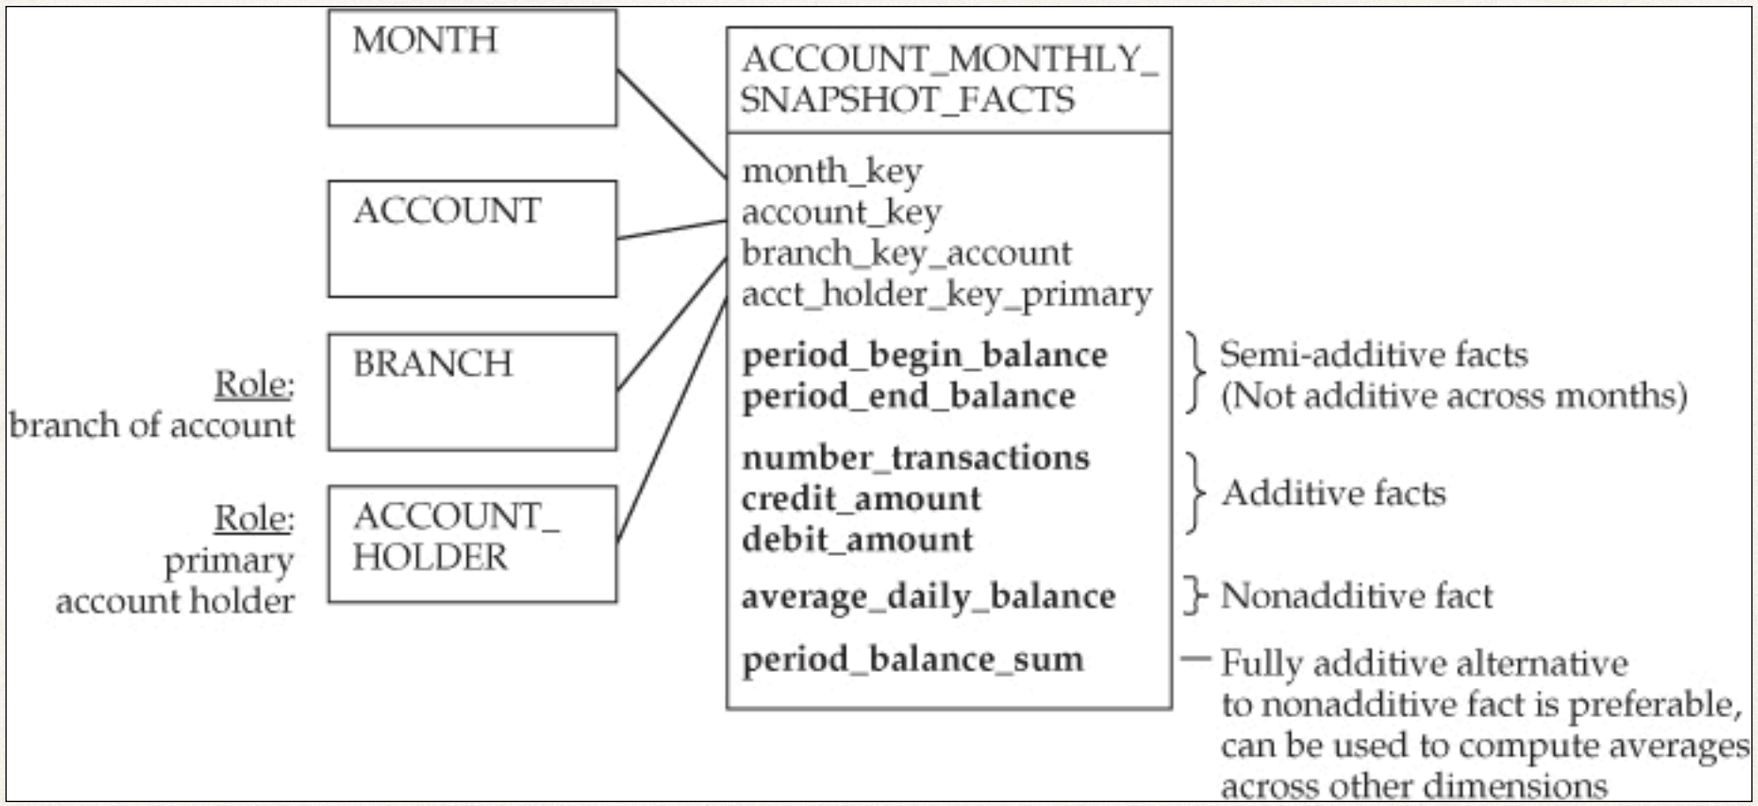

Fig. 11-4

Additional snapshot facts fit well together - natural affinity

Snapshot fact tables - differences

- **Dense**, not sparse
- **Periodic**, not event-based
- **Semi-additive** (e.g. account balances, temperatures)
- Granularity aligns with dimensions, not events

Transactions and Snapshots

- These two designs **complement** each other
- Each may be viable for different kinds of analysis
- To simplify ETL, transaction table should be source for snapshot tables
- In this way, snapshot tables **derive** from snapshots

Accumulating Snapshots

- Transaction fact tables support measuring simple processes (based on single events)
- Snapshot fact tables support measuring patterns in state over time (based on periodic observations)
- How do we measure complex processes?

What's a complex process?

- think "workflow" - more than one service station, more than one role type, and transitions between each

- measure time at each step in the process - **intervals**, not events

- use to find **bottlenecks** to feed into process modeling and simulation, among others

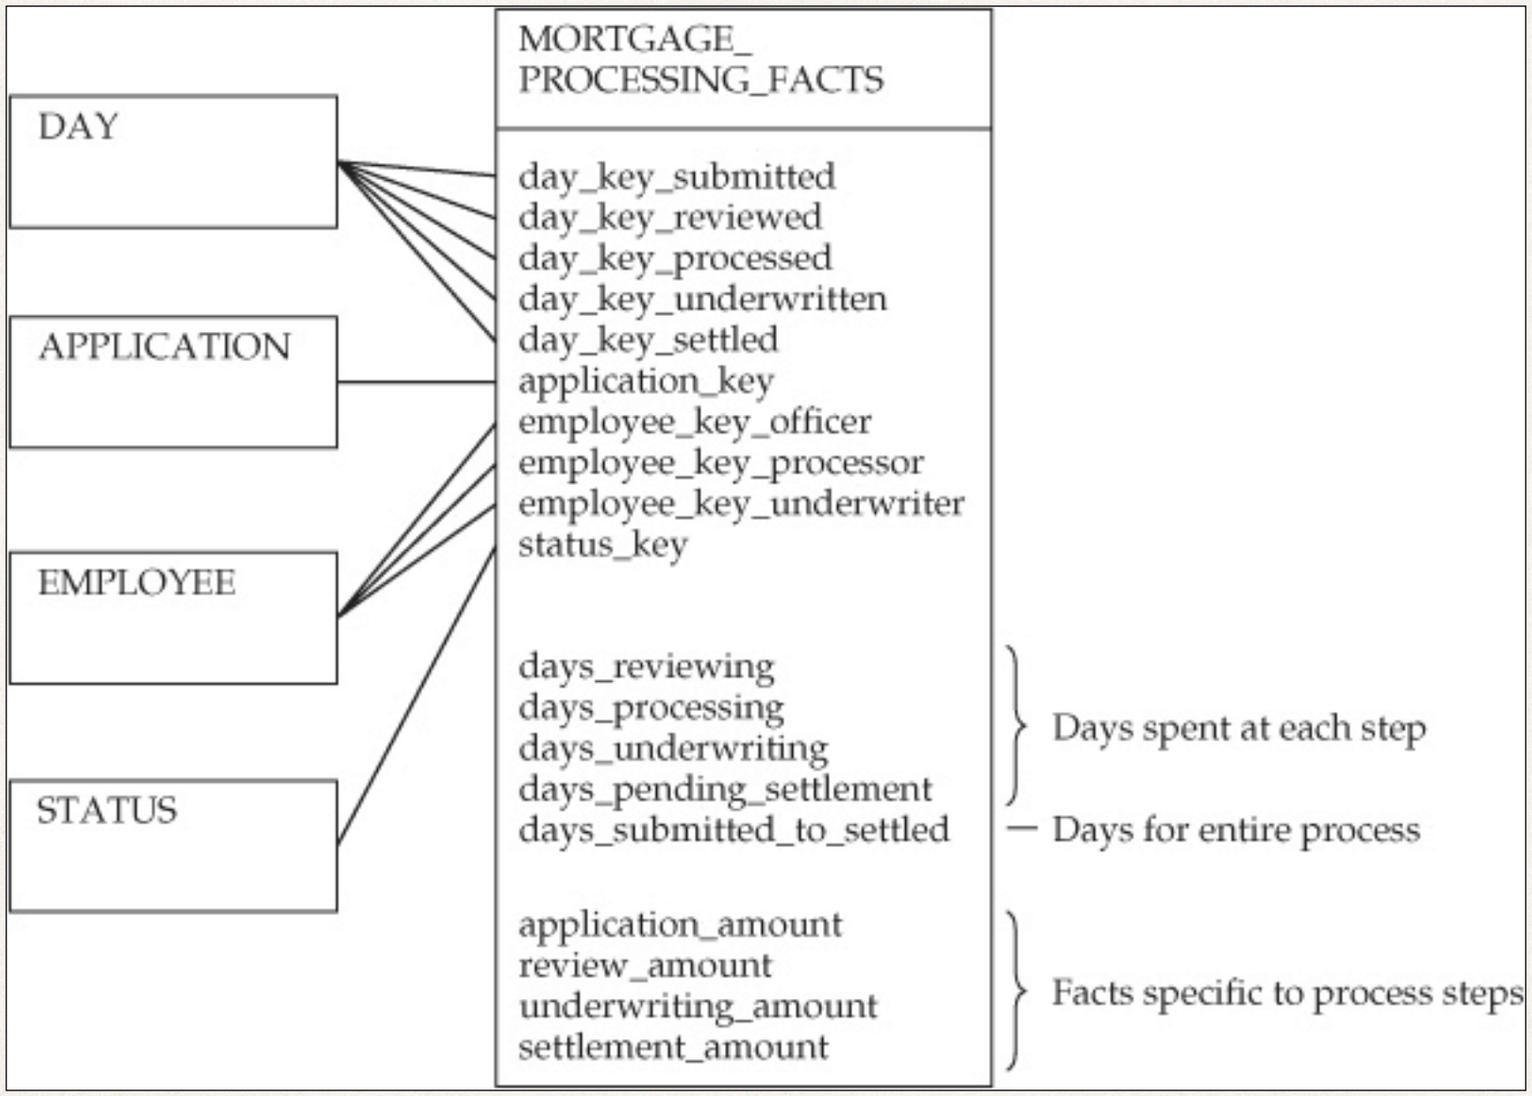

Fig. 11-6

Accumulating snapshot fact table design

Values accumulate over time

- Each row represents one process - multiple steps
- Row is updated as time passes and processing continues at one stage or another
- Time may pass with little or no progress; time elapsed increases at one stage or another

On Day 1 (Submitted; under review by officer):

application_ key	day_key_ submitted	day_key_ reviewed	day_key_ processed	day_key_ underwritten	day_key_ closing	application_ amount	review_ amount	underwriting_ amount	days_ reviewing	days_ processing	...
1011	**1021**	0000	0000	0000	0000	**100,000**	0	0	0	0	

Day 2 (No status change):

application_ key	day_key_ submitted	day_key_ reviewed	day_key_ processed	day_key_ underwritten	day_key_ closing	application_ amount	review_ amount	underwriting_ amount	days_ reviewing	days_ processing	...
1011	1021	0000	0000	0000	0000	100,000	0	0	**1**	0	

Days 3–9 (not shown)...

Day 10 (Reviewed; documents being gathered by processor):

application_ key	day_key_ submitted	day_key_ reviewed	day_key_ processed	day_key_ underwritten	day_key_ closing	application_ amount	review_ amount	underwriting_ amount	days_ reviewing	days_ processing	...
1011	1021	**1031**	0000	0000	0000	100,000	**90,000**	0	**9**	0	

Day 11 (No status change):

application_ key	day_key_ submitted	day_key_ reviewed	day_key_ processed	day_key_ underwritten	day_key_ closing	application_ amount	review_ amount	underwriting_ amount	days_ reviewing	days_ processing	...
1011	1021	1031	0000	0000	0000	100,000	90,000	0	9	**1**	

Remaining steps...

Fig. 11-7

Evolution of one accumulating snapshot record

Accumulating vs. transaction facts

- Accumulating snapshots excel at capturing / updating intervals

- Ideal for studying performance for each stage

- Complements transaction fact tables where volume, workloads, etc. may be more easily queried

- Like snapshots, transaction fact table should be source to simplify ETL

Project 02 work session