

ISTM 6212 - Week 10

Map / Reduce, Trifecta Wrangler

Daniel Chudnov, dchud@gwu.edu

Agenda

- ❖ Schedule check
- ❖ Map / Reduce
- ❖ Project 03
- ❖ Wrangling with Trifecta
- ❖ Exercise 05 work session

Schedule check

Map / Reduce

How did we get here?

- ❖ In the beginning: the command line, pipes and filters
- ❖ Relational databases, normalization, transactions
- ❖ Dimensional models, denormalization, conformance
- ❖ What does each bring us?

Command line - benefits

- ❖ Fast, free, ubiquitous, flexible
- ❖ Small, focused tools do one thing well
- ❖ Text / line orientation keeps things simple
- ❖ Pipeline model allows sophisticated workflows
- ❖ Filters can be in any language
- ❖ Scales up pretty well, to an extent

Command line - drawbacks

- ❖ Feels a bit arcane
- ❖ Hundreds of little things to learn
- ❖ Easy to screw up
- ❖ Different environments vary (e.g. Ubuntu vs. MacOS)
- ❖ Takes a lot of practice, helps to know a wizard

Relational databases - benefits

- ❖ Rock-solid basis in relational algebra
- ❖ Wealth of commercial and free / open source products
- ❖ Long-developed SQL specs for DDL, DML
- ❖ Sophisticated tools for access control, scaling up, data consistency, application development
- ❖ SQL mostly the same across systems

Relational databases - drawbacks

- ❖ SQL mostly the same across systems
- ❖ Heavy-duty products require dedicated staff
- ❖ Overhead of schema development, normalization
- ❖ Schema requirements limit flexibility, slow development
- ❖ Transactional designs not ideal for analysis

Dimensional models - benefits

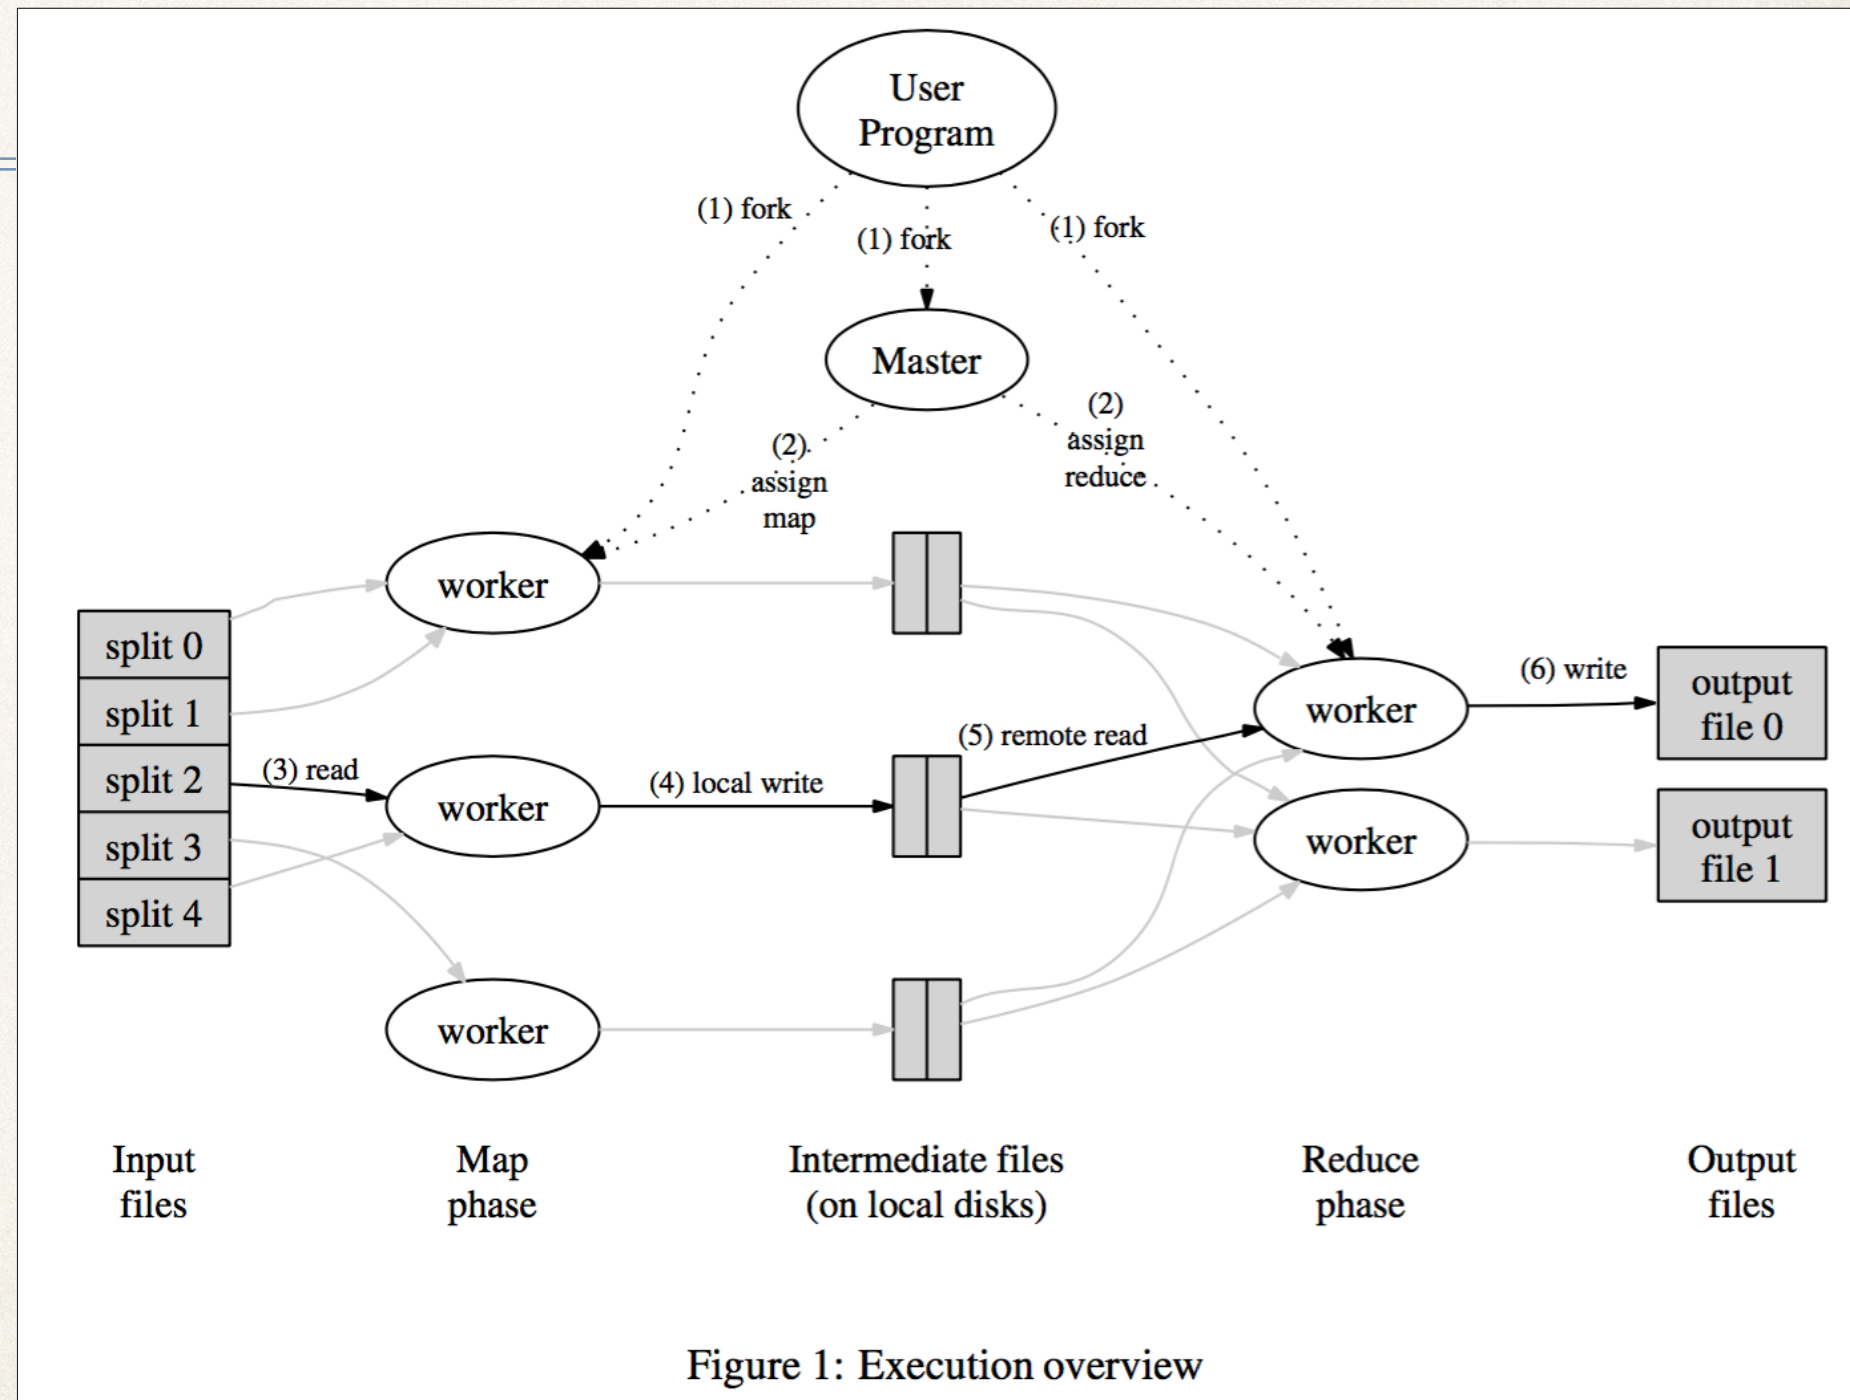
- ❖ Time-tested concepts: dimension, fact, grain, conformance
- ❖ Optimal strategies for analytical needs
- ❖ Allow analysts to move quickly in extracting data
- ❖ Denormalization keeps schema and queries relatively simple
- ❖ BI tool integration simplifies focus on analysis, reporting

Dimensional models - drawbacks

- ❖ Still requires dedicated staff
- ❖ Still requires schema overhead
- ❖ ETL process must be carefully managed
- ❖ Ad hoc designs don't grow and scale well
- ❖ Centrally managed designs may move too slowly

Map / Reduce

- ❖ Dean and Ghemawat, 2004
- ❖ Simplified programming model
- ❖ Flexible data model
- ❖ Optimized computing model



Map / Reduce

- ❖ Original paradigm 50+ years old (functional programming in Lisp, etc.)
- ❖ Developed at Google in the early 2000s
- ❖ Concept spread widely, Apache Hadoop project 10+ years old (see en.wikipedia.org/wiki/Apache_Hadoop)
- ❖ Mature platform ~5 years, rapidly growing ecosystem

How does it work?

- ❖ Start with any source data
- ❖ **Map** function extracts keys and values
- ❖ **Reduce** function summarizes values by key
- ❖ That's it!

Map function example

ALPS AND SANCTUARIES OF PIEDMONT AND THE CANTON TICINO
by Samuel Butler

Author's Preface to First Edition

❖ Source data:

I should perhaps apologise for publishing a work which professes to deal with the sanctuaries of Piedmont, and saying so little about the most important of them all--the Sacro Monte of Varallo. My excuse must be, that I found it impossible to deal with Varallo without making my book too long. Varallo requires a work to itself; I must, therefore, hope to return to it on another occasion.

❖ Mapped data:

1 alps
1 and
1 sanctuaries
1 of
1 piedmont
1 and
1 the
1 canton
1 ticino

Map function

❖ `grep -oE '\w{2,}' example.txt | tr '[:upper:]' '[:lower:]'`
`| uniq -c`

Reduce function example

❖ Mapped data:

1 alps
1 and
1 sanctuaries
1 of
1 piedmont
1 and
1 the
1 canton
1 ticino

❖ Reduced data:

1 about
1 all
1 alps
3 and
1 another
1 apologise
1 author
1 be
1 book

Reduce function

❖ `sort | uniq -c`

that's Map / Reduce*

*or most of it, anyway

Unique value from each

- ❖ Command line - free, flexible, & powerful assembly of small pieces
- ❖ RDBMS - consistent data management, standardized declarative query, statistics-driven optimization
- ❖ Dimensional - optimized query experience for analysts
- ❖ Map / Reduce - flexible data, simplified code model, optimized computing strategy scales linearly

How map/reduce works

```
map(String key, String value):  
  // key: document name  
  // value: document contents  
  for each word w in value:  
    EmitIntermediate(w, "1");  
  
reduce(String key, Iterator values):  
  // key: a word  
  // values: a list of counts  
  int result = 0;  
  for each v in values:  
    result += ParseInt(v);  
  Emit(AsString(result));
```

- ❖ Simplified programming model
- ❖ Commodity hardware
- ❖ Scheduler manages jobs, storage
- ❖ Network file system helps jobs run near data

How map/reduce works (2)

- ❖ Input split into standard-sized chunks
- ❖ Chunked data stored redundantly on networked file system
- ❖ Scheduler assigned compute nodes to map data they have stored locally (“locality”)
- ❖ Mapped data written back to file system

How map/reduce works (3)

- ❖ Reduce jobs scheduled to work on mapped data
- ❖ Reducers sort intermediate keys to group / combine results
- ❖ Reducers write results back to file system
- ❖ Scheduler tracks storage, job status, re-assigns around failures, duplicates workload near the end

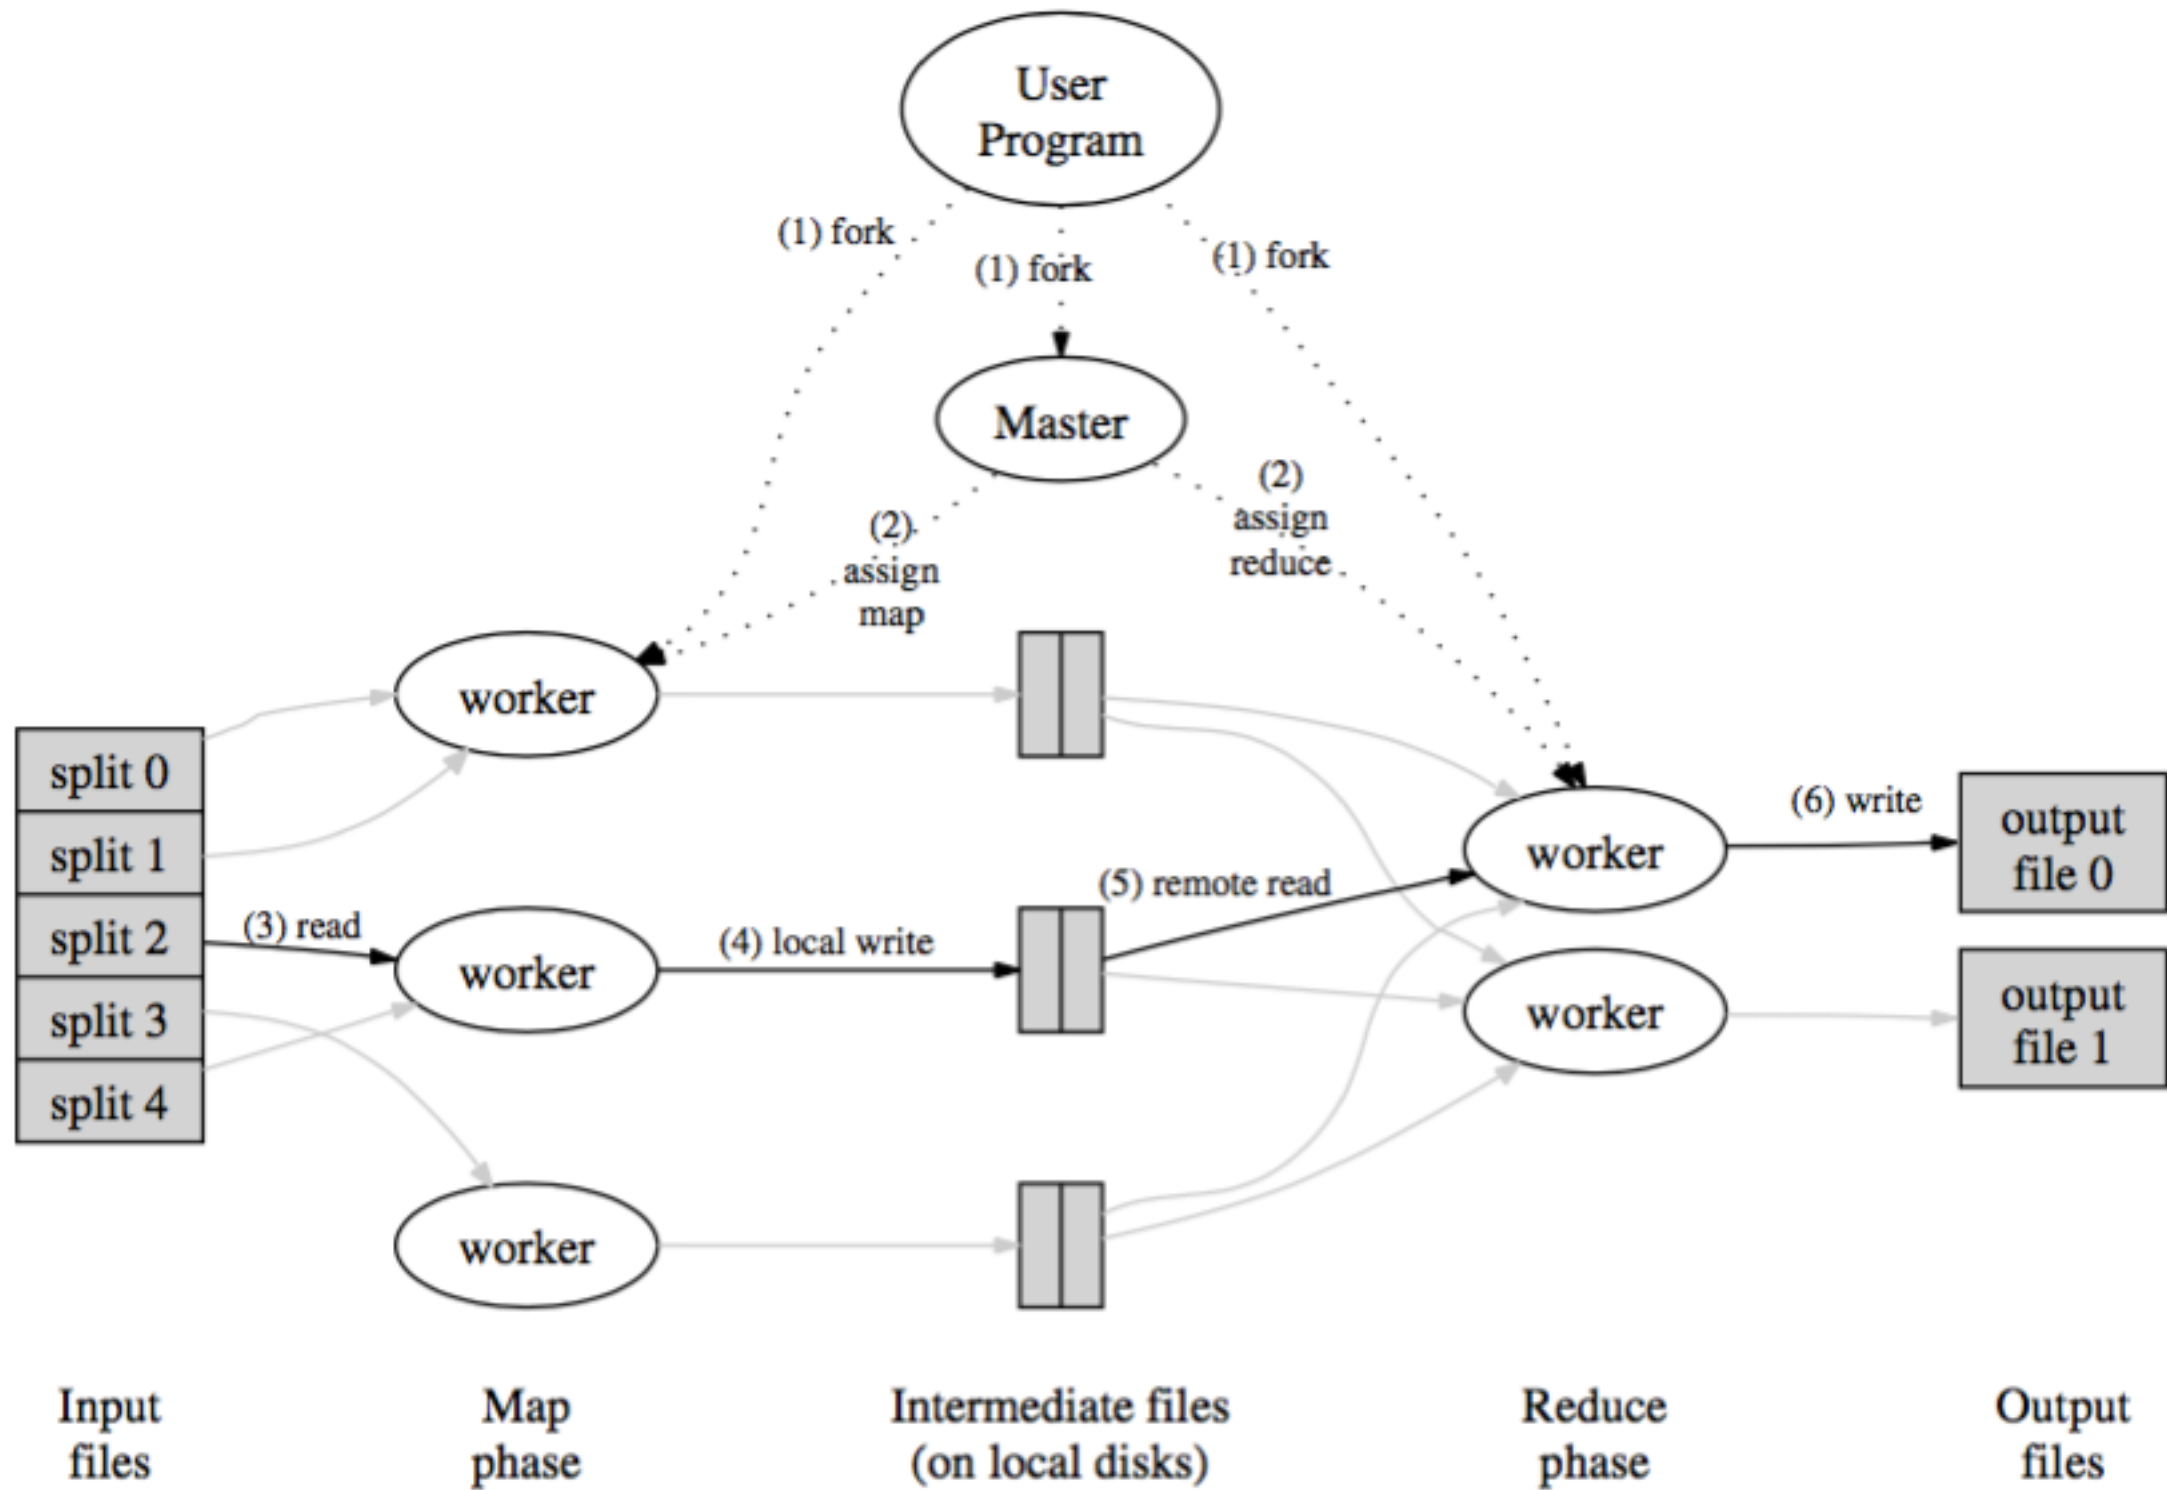




Figure 1: Execution overview

Compute optimizations


- ❖ **Locality** of redundant data ensures jobs occur on nodes with copies of data - no extra network traffic
- ❖ **Fault tolerance** allows jobs to finish even if whole blocks of machines become unavailable
- ❖ **Backup tasks** route around “stragglers”
- ❖ Hash optimization, bad records, counters, etc.



(a) Normal execution



(b) No backup tasks



(c) 200 tasks killed

Figure 3: Data transfer rates over time for different executions of the sort program

What does this mean for you?

- ❖ Map / reduce is widely used; **Hadoop** is a mature free / open source implementation
- ❖ Most organizations using data at scale have clusters onsite, via AWS or BigTable+Google, or plan to
- ❖ Schema-free flexibility, simple programming model, great performance means we can just “pile up our data” and transform / conform it when we need it

How is this good for you?

- ❖ **Data lakes*** still require conforming dimensions for merging / drilling across sources
- ❖ **Spark** and dozens of other tools build on and improve map / reduce model and Hadoop infrastructure layer
- ❖ You can run it on your laptop or via AWS / Google
- ❖ You can get started with it easily (Exercise 06 next week!)

*search “Tamara Dull data lake” for more on this concept

Is all this stuff mature?

- ❖ Yes, or at least mature enough
- ❖ You can download Spark and use it locally with minimal configuration
- ❖ Cloud-hosted services already stable and inexpensive
- ❖ SQL interfaces (Spark SQL, Apache Drill, etc.) are coming quickly and improving quickly
- ❖ SQL on top of map / reduce - SQL knowledge counts

More on overhead

- ❖ How big is your data chunk size?
- ❖ Which hashing algorithm fits best?
- ❖ Which compression should you use?
- ❖ How many nodes are needed per job?
- ❖ What level of storage redundancy makes sense?
- ❖ Where do dimension conformance specifications live?

Keep it in context

- ❖ Schema vs schemaless
- ❖ Consistency vs flexibility
- ❖ Speed vs scale
- ❖ Scaling requires staffing
- ❖ Scaling requires scheduling overhead

Next week: Spark

- ❖ skim the Apache Spark site, docs
- ❖ get a Databricks Community Edition account
- ❖ Exercise 06 - Spark on Databricks (probably!)

Project 03 (Final project)

Final project

- ❖ Groups of four people (should be 12!)
- ❖ Task like Project 02, but open to tools: commandline, db, spark, trifacta, etc.
- ❖ Dataset \geq 250,000 records
- ❖ Presentation Tuesday, 6 December
- ❖ Writeup Friday, 9 December
- ❖ Reviews Tuesday, 12 December

Trifacta Wrangler

Trifacta Wrangler

- ❖ Grew out of university research
- ❖ Well-funded, quickly improving
- ❖ Aims to simplify data wrangling, push ETL out to analysts
- ❖ Example workflow: Trifacta for cleaning / wrangling, Tableau for visualization

Trifacta Wrangler - concepts

- ❖ Projects and datasets
- ❖ Transform scripts
 - ❖ Suggested transform patterns
 - ❖ Required transform language
- ❖ Develop scripts of data samples
- ❖ Generate results - output to CSV, Excel, TDE

Exercise 05 work session
